



365Days

English Stydy Plan

Privacy in ML

Randomized Response

Truth x_1, \dots, x_n $x_i \in \{0, 1\}$

Revealed

$$y_i = \begin{cases} \underline{x_i} & \text{w.p. } \frac{e^\epsilon}{1+e^\epsilon} \quad \checkmark \quad \text{Truth} \\ 1-x_i & \text{w.p. } \frac{1}{1+e^\epsilon} \quad \checkmark \quad \text{Falsehood} \end{cases}$$

Consider two datasets

Truth $x = \{x_1, \dots, x_n\}$ $x' = \{x'_1, \dots, x'_n\}$

Assume $x_i = x'_i$ $\forall i = \{1, \dots, n-1\}$

Revealed \downarrow Randomized response
 $RR(x) = y = [y_1, \dots, y_n]$

$RR(x') = y' = [y'_1, \dots, y'_n]$

$$P\left(\underset{\substack{\uparrow \\ [b_1 \dots b_n]}}{RR(x)} = b\right) = P\left([y_1, y_2, \dots, y_n] = [b_1, b_2, \dots, b_n]\right)$$

$$= \left(\prod_{i=1}^{n-1} P(y_i = b_i) \right) \underbrace{P(y_n = b_n)}_{\rightarrow \textcircled{1}}$$

$$P\left(RR(x') = b\right) = \prod_{i=1}^n P(y_i' = b_i)$$

$$= \left(\prod_{i=1}^{n-1} P(y_i' = b_i) \right) P(y_n' = b_n)$$

$$\rightarrow \left(\prod_{i=1}^{n-1} P(y_i = b_i) \right) \underbrace{P(y_n' = b_n)}_{\rightarrow \textcircled{2}}$$

$$\frac{P(y_n = b_n)}{P(y_n' = b_n)} = \begin{cases} \frac{e^\epsilon / 1 + e^\epsilon}{1 / 1 + e^\epsilon} = e^\epsilon & \text{if } \underline{b_n = x_n} \\ \frac{1 / 1 + e^\epsilon}{e^\epsilon / 1 + e^\epsilon} = e^{-\epsilon} & \text{if } b_n = 1 - x_n \end{cases}$$

$$\frac{P(y_n = b_n)}{P(y_n' = b_n)} \leq e^\epsilon$$

$$\boxed{P(y_n = b_n) \leq e^\epsilon P(y_n' = b_n)} \quad \text{— substitute in ①.}$$

$$P(\mathcal{RR}(x) = b) = \left(\prod_{i=1}^{n-1} P(y_i = b_i) \right) \underbrace{P(y_n = b_n)}$$

$$\leq \left(\prod_{i=1}^{n-1} P(y_i' = b_i) \right) \underbrace{e^\epsilon P(y_n' = b_n)}$$

$$= e^\epsilon \underbrace{\prod_{i=1}^n P(y_i' = b_i)}$$

$$P(\mathcal{RR}(x') = b)$$

\Rightarrow

$$\frac{P(RR(x)=b)}{P(RR(x')=b)} \leq e^\epsilon$$

How good is this mechanism

want

$$\frac{1}{n} \sum_{i=1}^n x_i$$

x_i
 $\{0,1\}$

Randomization




x_i
 $E[x_i]$

$$E[y_i] = \left(\frac{e^\epsilon}{1+e^\epsilon} \right) x_i + \left(\frac{1}{1+e^\epsilon} \right) (1-x_i)$$

$$\frac{e^\epsilon x_i + 1 - x_i}{1 + e^\epsilon}$$

$$E[y_i] = \frac{x_i \left(\frac{e^\epsilon - 1}{e^\epsilon + 1} \right) + \frac{1}{1+e^\epsilon}}{1}$$



$$z_i = \left(y_i - \frac{1}{1+e^\epsilon} \right) \left(\frac{e^\epsilon + 1}{e^\epsilon - 1} \right)$$

$E[z_i] = ?$ x_i (exercise)

x_1, \dots, x_n

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

\downarrow

\downarrow

y_1, \dots, y_n

y_n

\downarrow

\downarrow

z_1

z_n

Guess for \bar{x} is $\bar{z} = \frac{1}{n} \sum_{i=1}^n z_i$

$$\mathbb{E}[\bar{z}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n z_i\right]$$

$$= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[z_i]$$

$$= \frac{1}{n} \sum_{i=1}^n x_i = \bar{x}$$

$$\left| \underbrace{\frac{1}{n} \sum_{i=1}^n x_i} - \underbrace{\frac{1}{n} \sum_{i=1}^n z_i} \right|$$

with high probability

$$\leq O\left(\frac{1}{\epsilon \sqrt{n}}\right) \leftarrow \underline{\text{utility}}$$

↑

TRUSTED CURATOR MODEL

Cynthia Dwork
[2010]

Differential Privacy

Let $M: \mathcal{X}^n \rightarrow \mathcal{Y}$. Consider two
"neighbouring" datasets $x, x' \in \mathcal{X}^n$

M is ϵ -D.P if for all x, x' neighbouring
and all $S \subseteq \mathcal{Y}$,

$$\frac{\Pr(M(x) \in S)}{\Pr(M(x') \in S)} \leq \begin{pmatrix} \epsilon \\ e \end{pmatrix}$$

LAPLACE MECHANISM

Sensitivity

$$f: X^n \rightarrow \mathbb{R} \quad (\text{average})$$

$$\Delta = \max_{\substack{\text{neighbouring} \\ \text{datasets} \\ x, x'}} |f(x) - f(x')|$$

$$\left| \frac{1}{n} \sum_{i=1}^n x_i - \frac{1}{n} \sum_{i=1}^n x'_i \right|$$



$$\left| \frac{1}{n} \sum_{i=1}^{n-1} x_i + \frac{1}{n} x_n - \frac{1}{n} \sum_{i=1}^{n-1} x_i - \frac{1}{n} x'_n \right|$$

$$= \frac{1}{n} |x_n - x'_n|$$

$$\Delta = \frac{1}{n}$$

Laplace distribution

$$f_{\text{Lap}}(x) = \frac{1}{2b} e^{-\frac{|x-\mu|}{b}}$$



$$\hat{x} = \frac{1}{n} \sum_{i=1}^n x_i + \eta$$

Laplace $(0, \frac{\Delta}{\epsilon})$


$$\mathbb{E}[\hat{x}] = \bar{x}$$

Can argue Laplace mechanism is

ϵ -DP





[Exercise].



utility

$$P \left(\left| \frac{1}{n} \sum_{i=1}^n x_i - \left(\frac{1}{n} \sum_{i=1}^n x_i + \eta \right) \right| \geq \epsilon \right)$$

$$\left. \begin{array}{c} \text{LAP} \\ O \left(\frac{1}{\epsilon n} \right) \end{array} \right| \begin{array}{c} \text{RR} \\ O \left(\frac{1}{\epsilon \sqrt{n}} \right) \end{array}$$



Approximate DP

$M: \mathcal{X}^n \rightarrow \mathcal{Y}$ is (ϵ, δ) D.P

if \forall neighbouring $x, x' \in \mathcal{X}^n$

and all $S \subseteq \mathcal{Y}$

$$P(M(x) \in S) \leq e^\epsilon P(M(x') \in S) + \delta$$

Add Gaussian noise

$$N\left(0, \ln\left(\frac{1}{\delta}\right) \frac{\Delta_2^2}{\epsilon^2}\right)$$

L2 Sensitivity

privacy

Laplacian noise $\text{Lap}\left(0, \frac{\Delta}{\epsilon}\right)$ $\left[\frac{d}{n\epsilon}\right]$

Gaussian noise $\mathcal{N}\left(0, \frac{\sqrt{d \log(1/\delta)}}{n\epsilon}\right)$

Why does it work?

$L_1 - \Delta$ $\max_{x, x'} \|f(x) - f(x')\|_1$

$L_2 - \Delta$ $\max_{x, x'} \|f(x) - f(x')\|_2$

$$\forall x \quad \|x\|_2 \leq \|x_1\|_2 \leq \sqrt{d} \|x_2\|$$

Properties of ADP

- Post-processing

- M is (ϵ, δ) DP

$F \circ M$ is also (ϵ, δ) DP


Composition

- $\{M_1, M_2, \dots, M_k\}$ are all (ϵ, δ) DP.

Basic: $(k\epsilon, k\delta)$

• Advanced:

$\left(\epsilon \sqrt{k \log(1/\delta')}, k\delta + \delta' \right)$ - DP
 $+ \epsilon(p^\epsilon - 1)$




Buyer = n

Valuation v_1, \dots, v_n



How to set price

$$p \rightarrow p \mid \{i : v_i \geq p\}$$



10, 10, 500

$S(x, p)$
revenue



$\chi \in \chi^n$ (valuations) (10, 10, 500)

\mathcal{H} (Prices)

$$S: (\chi^n \times \mathcal{H}) \rightarrow \mathbb{R}$$

$$\Delta = \max_{h \in \mathcal{H}} \max_{x, x'} |S(x, h) - S(x', h)|$$

EXPONENTIAL MECHANISM (2018 Taiwan et.al).

Select $h \in \mathcal{H}$ with probability

proportional to

$$\frac{e^{S(x, h)}}{e^{\frac{2\Delta}{\epsilon}}}$$

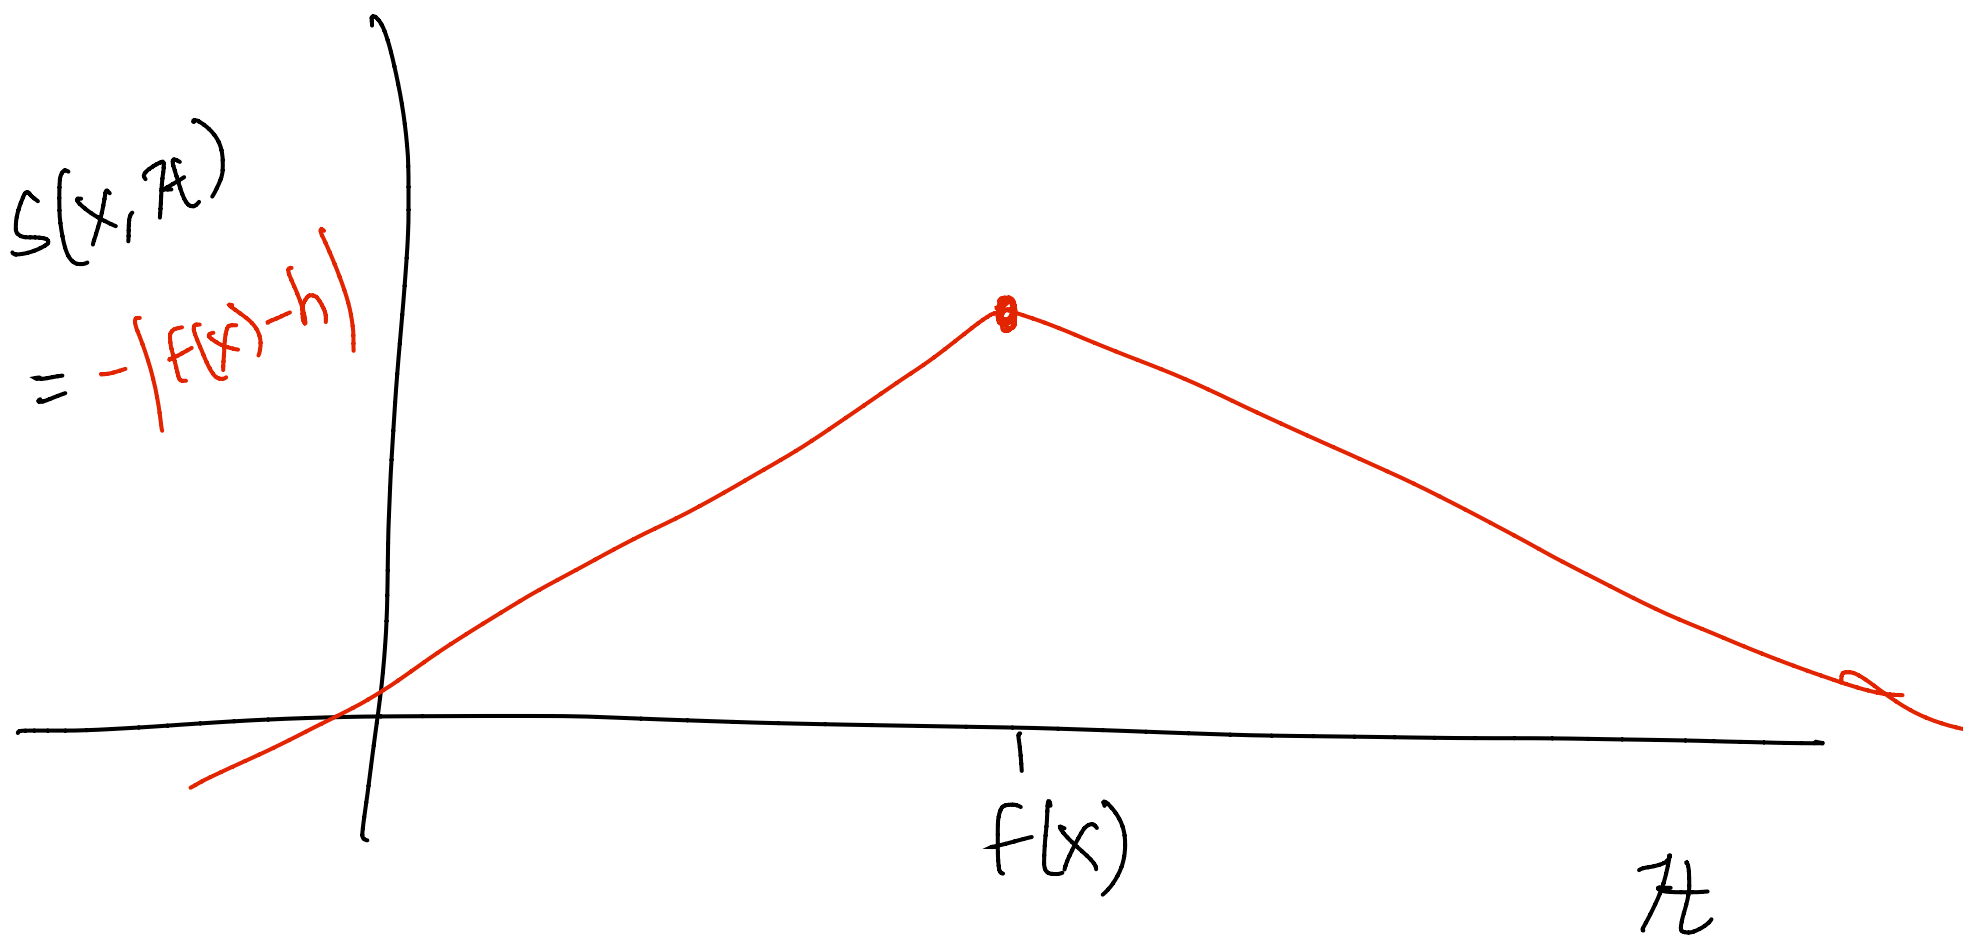
EM is ϵ -DP.

Utility

valuations

$$P \left(\underbrace{S(\underbrace{EM(x)}_{\text{Price}})}_{\text{Revenue}} \leq \underbrace{OPT(x)}_{\substack{\text{max} \\ \text{money} \\ \text{I} \\ \text{can} \\ \text{make.}}} \left(\frac{2\Delta}{\epsilon} \ln(1/\delta) + t \right) \right) \leq e^{-t}$$

Laplace Mechanism



Privacy in ML

$$L(w) = \frac{1}{n} \sum_{i=1}^n \underbrace{\ell(x_i, y_i; w)}_{\text{Loss function}} + \frac{R(w)}{n}$$

$$w^* = \arg \min_w L(w)$$

- Output perturbation

$$\hat{w} = w^* + \eta \leftarrow \text{noise. } \eta \in \mathbb{R}^d.$$

$$(\epsilon, \delta) \stackrel{\text{Excess error}}{O} \left(\frac{d}{\epsilon \sqrt{n}} \right)$$

- Objective perturbation

$$\hat{w} = \arg \min_w \left(L(w) + \underbrace{\hat{w}^T \eta}_{\text{noise.}} \right)$$

(ϵ, δ) -D.P

$$O\left(\frac{d}{\epsilon\sqrt{n}}\right)$$

Gradient perturbation

Run SGD with "noisy gradients".

(ϵ, δ)

$$O\left(\frac{\sqrt{d}}{\epsilon n}\right)$$
