

Introduction to ML

ACM India Summer School on
Responsible & Safe AI

3rd June, 2024



There is a function f in nature and we want to model / approximate it:

$$f : D \rightarrow R$$

We only have input - output samples available

Example:

- ▶ D : Number Tuples
- ▶ R : Real numbers

$[(1,2), 3], [(2,3), 5], [(3,4), 7], [(4,5), 9]$

Game: Guess the function f

Parameters and architecture

To make an approximation, we *parameterize* the function. We call this function f_θ

$$f_\theta : D \rightarrow R$$

Continuing with our previous example, let us define f_θ as:

$$f_\theta(x) = \theta_2 x_1 + \theta_1 x_0 + \theta_0$$

We have effectively defined an *architecture* for our function with parameters. Now, our problem becomes finding the right values for $\theta_0, \theta_1, \theta_2$ that “fits” the data.

There are non – parametric Machine Learning Algorithms as well! E.g : knn, Decision Trees.

ML Pipeline

Data

Representation

Algorithms

Evaluation

Data Preprocessing – Textual data

- Removing punctuations like . , ! \$ () * % @
- Removing URLs
- Removing Stop words – the, it, a, was
- Lower casing
- Tokenization – break a sentence into small words
- Stemming vs Lemmatization



Data Preprocessing – Textual data

Stemming	Lemmatization
Stemming is a process that stems or removes last few characters from a word, often leading to incorrect meanings and spelling.	Lemmatization considers the context and converts the word to its meaningful base form, which is called Lemma.
For instance, stemming the word ' Caring ' would return ' Car '.	For instance, lemmatizing the word ' Caring ' would return ' Care '.
Stemming is used in case of large dataset where performance is an issue.	Lemmatization is computationally expensive since it involves look-up tables and what not.

Data Representation

Bag of Words Example

Document 1

The quick brown fox jumped over the lazy dog's back.

Document 2

Now is the time for all good men to come to the aid of their party.

Term	Document 1	Document 2
aid	0	1
all	0	1
back	1	0
brown	1	0
come	0	1
dog	1	0
fox	1	0
good	0	1
jump	1	0
lazy	1	0
men	0	1
now	0	1
over	1	0
party	0	1
quick	1	0
their	0	1
time	0	1

Stopword List

for
is
of
the
to

TF-IDF Example

Word	TF		IDF	TF*IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2) = 0$	0	0
Car	1/7	0	$\log(2/1) = 0.3$	0.043	0
Truck	0	1/7	$\log(2/1) = 0.3$	0	0.043
Is	1/7	1/7	$\log(2/2) = 0$	0	0
Driven	1/7	1/7	$\log(2/2) = 0$	0	0
On	1/7	1/7	$\log(2/2) = 0$	0	0
The	1/7	1/7	$\log(2/2) = 0$	0	0
Road	1/7	0	$\log(2/1) = 0.3$	0.043	0
Highway	0	1/7	$\log(2/1) = 0.3$	0	0.043

Data Representation

Bag of Words Example

Document 1

The quick brown fox jumped over the lazy dog's back.

Document 2

Now is the time for all good men to come to the aid of their party.

Term	Document 1	Document 2
aid	0	1
all	0	1
back	1	0
brown	1	0
come	0	1
dog	1	0
fox	1	0
good	0	1
jump	1	0
lazy	1	0
men	0	1
now	0	1
over	1	0
party	0	1
quick	1	0
their	0	1
time	0	1

Stopword List

for
is
of
the
to

TF-IDF Example

Word	TF		IDF	TF*IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2) = 0$	0	0
Car	1/7	0	$\log(2/1) = 0.3$	0.043	0
Truck	0	1/7	$\log(2/1) = 0.3$	0	0.043
Is	1/7	1/7	$\log(2/2) = 0$	0	0
Driven	1/7	1/7	$\log(2/2) = 0$	0	0
On	1/7	1/7	$\log(2/2) = 0$	0	0
The	1/7	1/7	$\log(2/2) = 0$	0	0
Road	1/7	0	$\log(2/1) = 0.3$	0.043	0
Highway	0	1/7	$\log(2/1) = 0.3$	0	0.043

Can be extended to n-grams.

Instead of counting words – counts grouping of words

Unigram - "The", "Car", "Truck"....

Bigram - "<s> The", "The Car", "Car Truck"....

Trigram - "<s> The Car", "The Car Truck" ...

Representing words by their context



- **Distributional semantics:** A word's meaning is given by the words that frequently appear close-by
 - “*You shall know a word by the company it keeps*” (J. R. Firth 1957: 11)
 - One of the most successful ideas of modern statistical NLP!
- When a word w appears in a text, its **context** is the set of words that appear nearby (within a fixed-size window).
- We use the many contexts of w to build up a representation of w

...government debt problems turning into **banking** crises as happened in 2009...

...saying that Europe needs unified **banking** regulation to replace the hodgepodge...

...India has just given its **banking** system a shot in the arm...

These **context words** will represent **banking**

Word vectors

We will build a dense vector for each word, chosen so that it is similar to vectors of words that appear in similar contexts, measuring similarity as the vector **dot** (scalar) **product**

$$\textit{banking} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

$$\textit{monetary} = \begin{pmatrix} 0.413 \\ 0.582 \\ -0.007 \\ 0.247 \\ 0.216 \\ -0.718 \\ 0.147 \\ 0.051 \end{pmatrix}$$


Note: **word vectors** are also called **(word) embeddings** or **(neural) word representations**
They are a **distributed** representation

Word2vec: Overview

Word2vec is a framework for learning word vectors
(Mikolov et al. 2013)


Idea:

- We have a large corpus (“body”) of text: a long list of words
- Every word in a fixed vocabulary is represented by a **vector**
- Go through each position t in the text, which has a center word c and context (“outside”) words o
- Use the **similarity of the word vectors** for c and o to **calculate the probability** of o given c (or vice versa)
- **Keep adjusting the word vectors** to maximize this probability




Skip-gram model
(Mikolov et al. 2013)

ML Algorithms



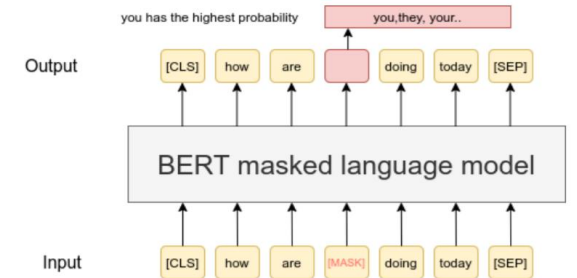
Supervised



Un-Supervised



Reinforcement



Self-Supervised

Supervised Learning

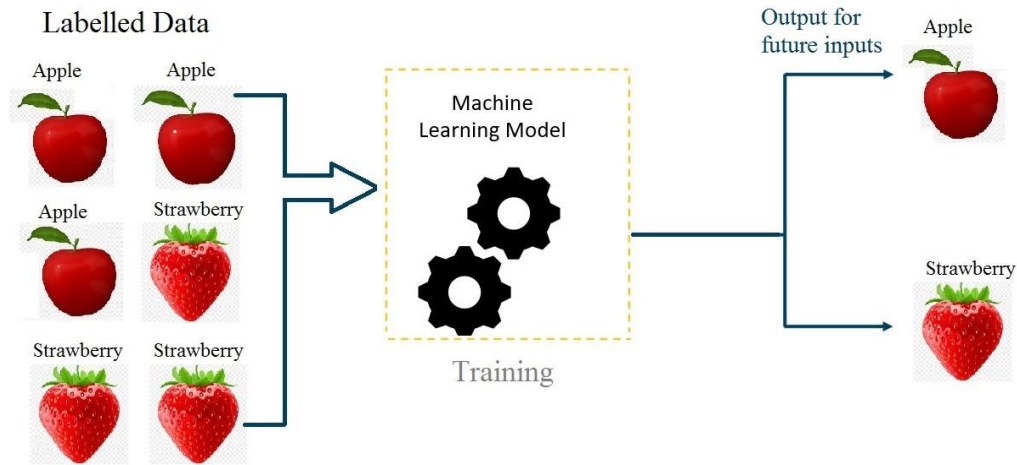
- Training on labeled dataset
- Learn mapping from inputs to outputs in training
- Predict output for new, unseen data

Given a dataset $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$, where x_i are the input features and y_i are the corresponding labels, the model learns a function f such that $f(x_i) \approx y_i$.

For regression: $y = f(x) + \epsilon$, where ϵ is the error term.

For classification: $P(y | x) = \frac{e^{f(x)}}{1 + e^{f(x)}}$ (logistic regression for binary classification).

Supervised Learning - Example



Classification: Determining if an email is spam or not spam



Regression: Predicting house prices based on features like size, location, and age

Supervised Learning

Common Algorithms

Linear Regression: $y = w^T x + b$

Logistic Regression:

$$P(y = 1 \mid x) = \frac{1}{1 + e^{-(w^T x + b)}}$$

Decision Trees: Recursive partitioning of data space.

Support Vector Machines (SVM): $y = \text{sign}(w^T x + b)$

Neural Networks: $y = \sigma(W^T x + b)$ (where σ is an activation function)

Un-Supervised Learning

- Train on data that does not have labeled responses
- The goal is to find hidden patterns in the input data

Given a dataset $\{x_1, x_2, \dots, x_n\}$, the objective is to find patterns or groupings in the data without explicit labels.

For clustering: Minimize the sum of squared distances between data points and their assigned cluster centroids,



$$\sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

where μ_i is the centroid of cluster C_i .

For dimensionality reduction (e.g., PCA): Maximize the variance retained in the projected data,

$$\max \|XW\|^2 \quad \text{subject to} \quad W^T W = I.$$

Unsupervised Learning - Example



Clustering: Grouping customers based on purchasing behavior.



Dimensionality Reduction: Reducing the number of features in a dataset while retaining its essential information (e.g., using PCA).

Metrics for Evaluation

Metrics for Evaluation

		Ground Truth	
Predictions		A	B
	A	98	0
	B	2	0

Confusion Matrix

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	true positive	false positive	precision = $\frac{tp}{tp+fp}$
	system negative	false negative	true negative	
		recall = $\frac{tp}{tp+fn}$	accuracy = $\frac{tp+tn}{tp+fp+tn+fn}$	

Confusion Matrix


		<i>gold standard labels</i>	
		gold positive	gold negative
<i>system output labels</i>	system positive	true positive	false positive
	system negative	false negative	true negative

Confusion Matrix

		<i>gold labels</i>			
		urgent	normal	spam	
<i>system output</i>	urgent	8	10	1	precision_u = $\frac{8}{8+10+1}$
	normal	5	60	50	precision_n = $\frac{60}{5+60+50}$
	spam	3	30	200	precision_s = $\frac{200}{3+30+200}$
		recall_u = $\frac{8}{8+5+3}$	recall_n = $\frac{60}{10+60+30}$	recall_s = $\frac{200}{1+50+200}$	

Let's calculate other metrics

- Precision = $2/(2+1) = 67\%$
- Recall = $2/(2+3) = 40\%$
- Accuracy = $(94+2)/100 = 96\%$
- Which measure should we account?
- RECALL





Another example


- A system which needs to launch a missile at a terrorist hideout located in a dense urban area.
- Which metric should we consider?
- Precision not 100% - civilian casualties

Accuracy vs Precision vs Recall


- Monitor Precision if a false positive carries higher cost.
- Monitor Recall if a false negative carries higher cost.



% of correct predictions



% of + class correctly predicted
[aka Recall / TPR]



correct prediction of + class
[aka Precision]

F1- Score

- What to do when one classifier has better precision but worse Recall, while other classifier behaves exactly opposite?

$$\begin{aligned} \text{F1 Score} &= \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} \\ &= \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \end{aligned}$$

- F1 measure punishes extreme values more !
- Definition of Recall and Precision have same numerator, different denominators. A sensible way to combine them is harmonic mean.

Evaluation metrics

- For Regression algorithms, we can't use accuracy as metric to judge models' performance
- We use Root Mean Square Error | Mean Absolute Error, among others

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}|$$

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2}$$


$$R^2 = 1 - \frac{\sum (y_i - \hat{y})^2}{\sum (y_i - \bar{y})^2}$$

Where,

\hat{y} – predicted value of y
 \bar{y} – mean value of y

Underfitting - Overfitting

- Underfitting: Happens when model is too simple to capture the underlying structure of the data
- Overfitting: This happens when a model is too complex and captures noise or random fluctuations in the training data



Bias- Variance


- **Bias (Underfitting)**

- Poor predictive performance as the model cannot capture the complexity of the problem.
- High bias leads to underfitting, where the model oversimplifies underlying patterns - Error due to overly simplistic assumptions in the learning algorithm.
- Example: Linear regression model assuming a strictly linear relationship for predicting housing prices.



- **Variance (Overfitting)**

- Poor generalization to unseen data due to essentially memorizing the training examples. Error due to excessive complexity in the learning algorithm.
- High variance captures both underlying patterns and noise in the training data.
- Example: Decision tree with a very deep structure trained on handwritten digits.

Bias- Variance Trade off



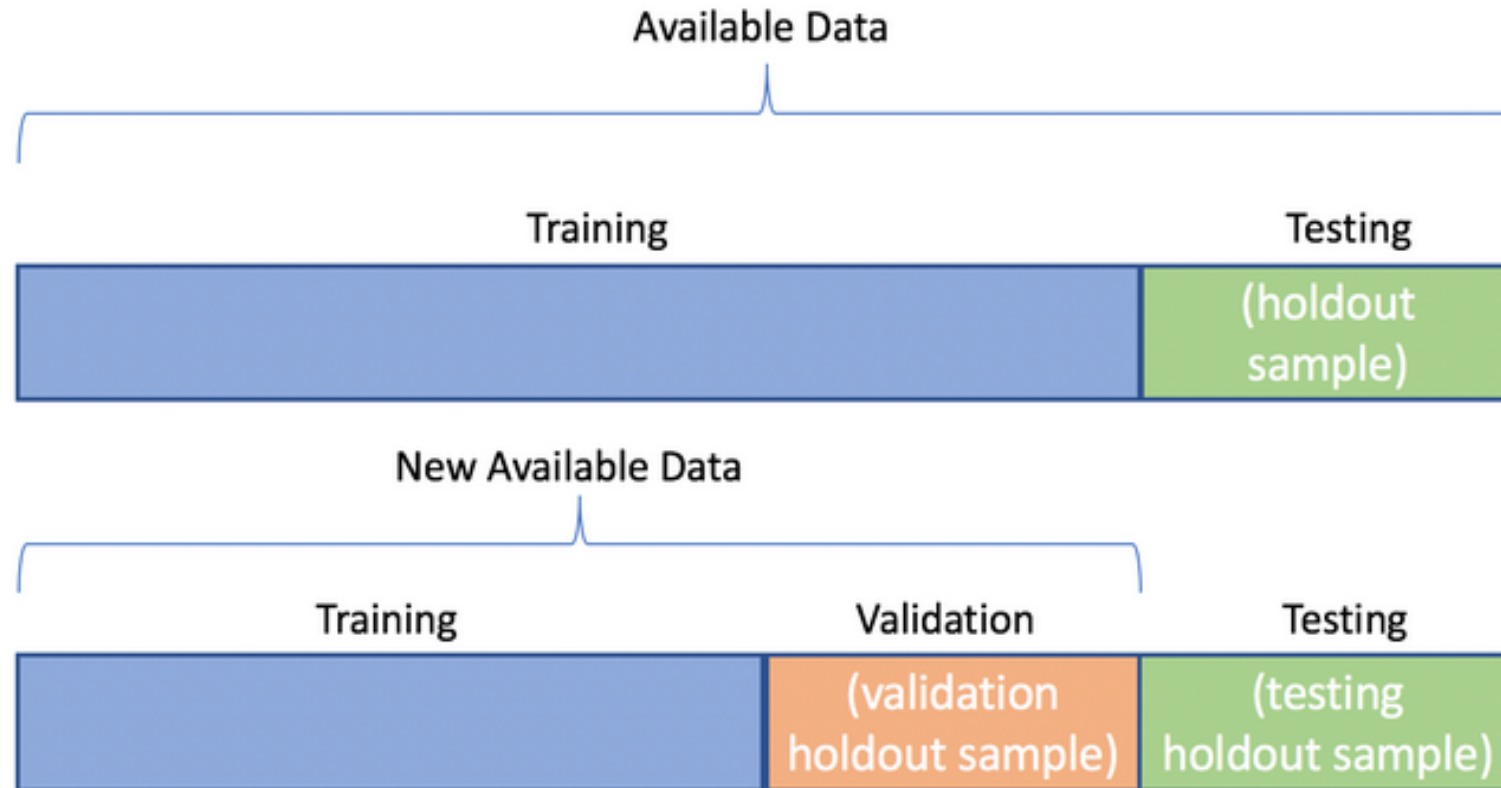
Bias- Variance Trade off



Bias- Variance Trade off

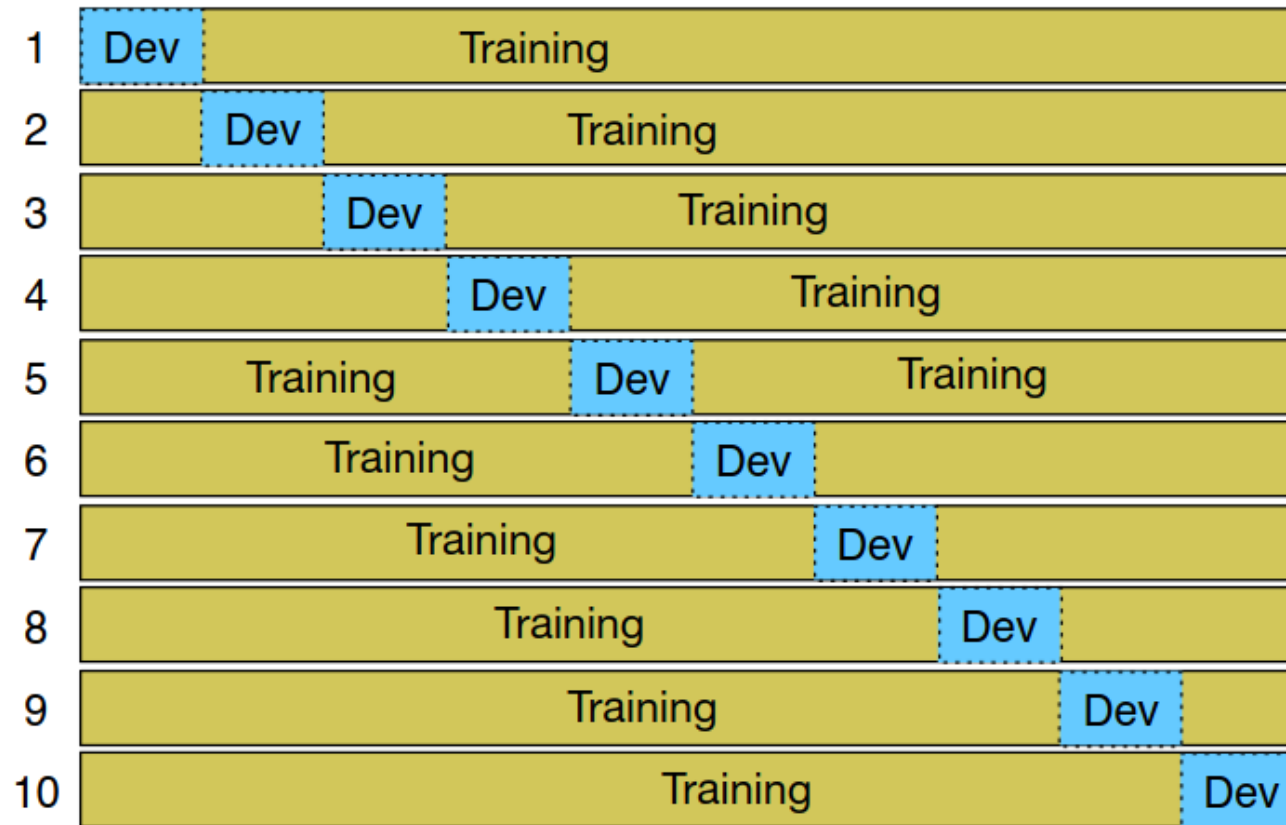
- The bias-variance tradeoff is the delicate equilibrium between underfitting and overfitting.
- The goal is to find the optimal level of complexity that allows a model to generalize effectively to unseen data.

Train-Val-Test paradigm



K-fold Cross Validation

Training Iterations



Testing

Test Set

Loss Functions

Why Loss Functions

- Let's say I am on the top of mountain and need to climb down. How do I decide where to walk towards?
 - Look around to see all the possible paths
 - Reject the ones going up
 - Finally, take the path that I think has the most slope downhill
- This intuition that I just judged my decisions against? This is exactly what loss function provides

Loss Functions


- Loss functions are used to quantify how well or bad a model can reproduce the values of the training set.
- The appropriate loss function depends on the type of problems and the algorithm we use.
- Loss function = Cost function = Objective function = Error function
- Loss function does not want to measure the entire performance of the network against a validation/test dataset

Loss Functions

The loss function is used to **guide the training process** in order to find a set of parameters that reduce the value of the loss function.



Loss Function



Training Process


Stochastic gradient descent

- Find a set of parameters which make the loss as small as possible.
- Change parameters at a rate determined by the partial derivatives of the loss function:

$$\frac{\partial \mathcal{L}}{\partial w} \quad \frac{\partial \mathcal{L}}{\partial b}$$

Properties

- Minimum when the output of the network is equal to the ground truth data
- Increase value when the output differs from the ground truth



Introduction to SK-learn

- Code